

MLton Hacker Guide

Stephen Weeks

May 3, 2016

This document describes how to hack MLton, a whole-program optimizing compiler for the Standard ML programming language. The MLton homepage is <http://www.mlton.org/MLton/>. The document contains an overview of the source tree, a description of the programming style used in MLton, and delves into the bowels of the compiler and associated tools.

This document is very incomplete.

Contents

1	The sources	3
2	The basis library	5
2.0.1	How MLton builds the basis environment	6
2.0.2	Modifying the basis library	6
2.0.3	The <code>misc</code> directory	7
2.0.4	Dead-code elimination	7
3	The runtime	9
4	MLton	10
4.1	Sources	10
4.2	Compiler Overview	12
5	Notes	14
5.1	IntInf and Flattener	14
6	Todo	16

Chapter 1

The sources

This section is an overview of the sources to the compiler and all of the associated tools. Here is a brief description of each element of the root source directory. Throughout the rest of this document, we will use pathnames that are relative to the source directory.

basis-library

The basis library implementation.

benchmark

Code and tests used for benchmarking MLton, SML/NJ, and Moscow ML.

bin

Scripts for type checking the basis library, making rpms, running MLton, and running regression tests.

doc

Sources for the user guide, hacker guide, web site, announcements, README.

include

Include files needed for compiling C files generated by MLton.

lib

SML library code, which is used in `mlton`, `mlprof`, and `benchmark`. There are also many generally useful libraries.

Makefile

To make everything. This is only used when building rpms.

man

Manual pages for `mlton` and `mlprof`.

mllex

Lexer generator, taken and slightly modified from SML/NJ.

`mlprof`

Profiler.

`mlton`

Compiler.

`mlyacc`

Parser generator, taken and slightly modified from SML/NJ.

`regression`

Regression tests, about 150 SML files that are used to test the compiler.

`runtime`

Runtime system, which includes the garbage collector and C libraries used in the basis (including the GMP used for `IntInf`).

Chapter 2

The basis library

The basis library is implemented with about 12,000 lines of SML code. There is roughly one file for each signature and structure that the library specification defines. The files are grouped in directories in the same way that the corresponding modules are grouped in the basis library documentation. Here is an overview of the `basis-library` directory.

`arrays-and-vectors` `general` `integer` `io` `list` `posix` `real` `system` `text`
SML code for basis library modules.

`basis.sml`
Automatically constructed by `bin/check-basis`. Used to type check the basis library under SML/NJ.

`bind-basis`
A list of the files (in order) that define what is exported by the basis library.

`build-basis`
A list of the files (in order) used to construct the basis library.

`Makefile`
Only has a target to clean the directory.

`misc`
SML code that didn't fit anywhere else. In particular, the `Primitive` structure.

`mlton`
The `MLton` structure, which is not part of the standard basis library. For more details on what `MLton` provides, see the `MLton User Guide`.

`sml-nj`
The `SMLofNJ` and `Unsafe` structures, which are not part of the standard basis library.

`top-level`

Files describing the overloads, infixes, modules, types, and values that the basis library makes available to user programs.

2.0.1 How MLton builds the basis environment

The `forceBasisLibrary` function in `mlton/main/compile.sml` builds the basis environment that is used to compile user programs. Conceptually, the basis environment is constructed in two steps. First, all of the files in `build-basis` are concatenated together and evaluated to produce an environment E . Then, all of the files in `bind-basis` are concatenated and evaluated in environment E to produce a new environment E' , which is the top-level environment. Another way to view it is that every user program is prefixed by the following.

```
local
  <concatenate files in build-basis>
in
  <concatenate files in bind-basis>
end
```

This view is not strictly accurate because some of the files are not SML (they use the `_prim`, `_ffi`, and `_overload` syntaxes) and because SML does not allow local functor or signature declarations. Here is a description of the basis files that are not SML.

`misc/primitive.sml`

Defines the `Primitive` structure, which binds (via the `_prim` syntax) all of the primitives provided by the compiler that the basis library uses.

`mlton/syslog.sml`

Defines constants and FFI routines used to implement `MLton.Syslog`.

`posix/primitive.sml`

Defines the `PosixPrimitive` structure, which binds the constants and FFI routines used to implement the `Posix` structure.

`top-level/overloads.sml`

Defines the overloaded variables available at the top-level the `_overload` syntax: `_overload x : ty as y_0 and y_1 and ...`

2.0.2 Modifying the basis library

If you modify the basis library, you should first check that your modifications are type correct using the `bin/check-basis` script. Since this MLton does not have a proper typechecker, this script uses SML/NJ. First, it concatenates the files as described in Section 2.0.1 into one file, `basis.sml`. It also replaces the nonstandard syntax (`_prim`, etc.) and declares the toplevel types to match MLton's

(necessary since SML/NJ uses 31 bits while MLton uses 32). It then feeds `basis.sml` to SML/NJ. If there are no type errors, a message like the following will appear.

```
stdIn:12213.1-12213.14 Error: operator is not a function [tycon mismatch]
  operator: unit
  in expression:
    () ()
```

This error message is intentionally introduced by `check-basis` at the end of `basis.sml` to make it clear that SML/NJ reached the end of `basis.sml` and has hence type checked the entire basis.

Once you have a basis library that type checks, you need to create a new version of MLton that uses this library. MLton preprocess the basis library to create a `world.mlton` file that contains the basis environment. The `world.mlton` file is stored in the `lib` directory and is loaded by `mlton` when compiling a user program (see the `bin/mlton` script). To build a new `world.mlton`, run `make world` from within the sources directory.

2.0.3 The misc directory

`cleaner.sig`

Functions for register “cleaning” functions to be run at certain times, in particular at program exit. The `TextIO` module uses these cleaners to ensure that IO buffers are flushed upon exit.

`suffix.sml`

Code that is (conceptually) concatenated on to the end of every user program. It just calls `OS.Process.exit`. The `forceBasisLibrary` function ensures that `suffix.sml` is elaborated in an environment where the basis library `OS` structure is available.

`top-level-handler.sml`

This defines the top level exception handler that is installed (via a special compiler primitive) in the basis library, before any user code is run.

2.0.4 Dead-code elimination

In order to compile small programs rapidly and to cut down on executable size, `mlton` runs a pass of dead-code elimination (`mlton/core-ml/dead-code.sig`) to eliminate as much of the basis library as possible. The dead-code elimination algorithm used is not safe in general, and only works because the basis library implementation has special properties:

- it terminates
- it performs no I/O
- it doesn’t side-effect top-level variables

The dead code elimination simply includes the minimal set of declarations from the basis so that there are no free variables in the user program (or basis). Hence, if you do something like the following in the basis, it will break.

```
val r = ref 13
val _ = r := 14
```

The dead code elimination will remove the `val _ = ...` binding.

Chapter 3

The runtime

There are multiple, possibly orthogonal issues. Limit checks and garbage collections are a little overloaded in their roles, because they also support preemptive thread switching and interrupt handling. Forcing frontier to be 0 and hitting a limit check (even a zero byte limit check) will invoke the GC, which will switch to the pending thread.

Recall that a limit check with `bytes = 0` really means a check for `LIMIT_SLOP` bytes (currently `LIMIT_SLOP = 512`).

Chapter 4

MLton

This chapter describes the compiler proper, which is found in the `mlton` directory.

4.1 Sources

`ast`

Abstract syntax trees produced by the front end.

`atoms`

Common atomic pieces of syntax trees used throughout the compiler, like constants, primitives, variables, and types.

`backend`

The backend translates from the `Cps` IL to a machine independent IL called `Machine`. It decides data representations, stack frame layouts, and creates runtime system information like limit checks and bitmasks.

`call-main.sml`

A one-line file that is the last line of the compiler sources. It calls the main function.

`closure-convert`

The closure converter, which converts from `Sxml`, the higher-order simply-typed IL, to `Cps`, the first-order simply-typed IL.

`cm`

Support for SML/NJ-style compilation manager (CM) files.

`codegen`

Both the C and the native X86 code generator.

`control`

Compiler switches used throughout the rest of the compiler.

core-ml

The implicitly typed IL that results from defunctorization. Contains a pass of dead code elimination for eliminating basis library code. Also contains the pass that replaces constants defined by `_prim` with their values.

elaborate

The elaborator, which matches variable uses with bindings in the AST IL and defunctorizes to produce a **CoreML** program. It does not do type checking yet, but will someday.

front-end

The lexer and parser, which turn files into ASTs.

main

The two main structures in the compiler, one (**Main**) for handling all the command line switches and one (**Compile**) which is a high-level view of the the compiler passes, from front end to code generation.

Makefile

To make the compiler.

mlton.cm

An automatically generated file (`make mlton.cm`) that lists all of the files (in order) that make up the compiler.

mlton.sml

An automatically generated file (`make mlton.sml`) that contains all of the compiler sources concatenated together.

rcps

An experimental IL, similar to CPS, but with more expressive types for describing representations (hence the “r”). Not yet in use.

sources.cm

For compiling with SML/NJ.

ssa

Static-Single-Assignment form, the first-order simply-typed IL on which most optimization is performed. There are roughly 20 different optimization passes (some of which run several times).

type-inference

The type inference pass, which translates from **CoreML** to **Xml**.

xml

The **Xml** and **Sxml** intermediate languages. Also, the passes that monomorphise, do polvariance, and implement exceptions.

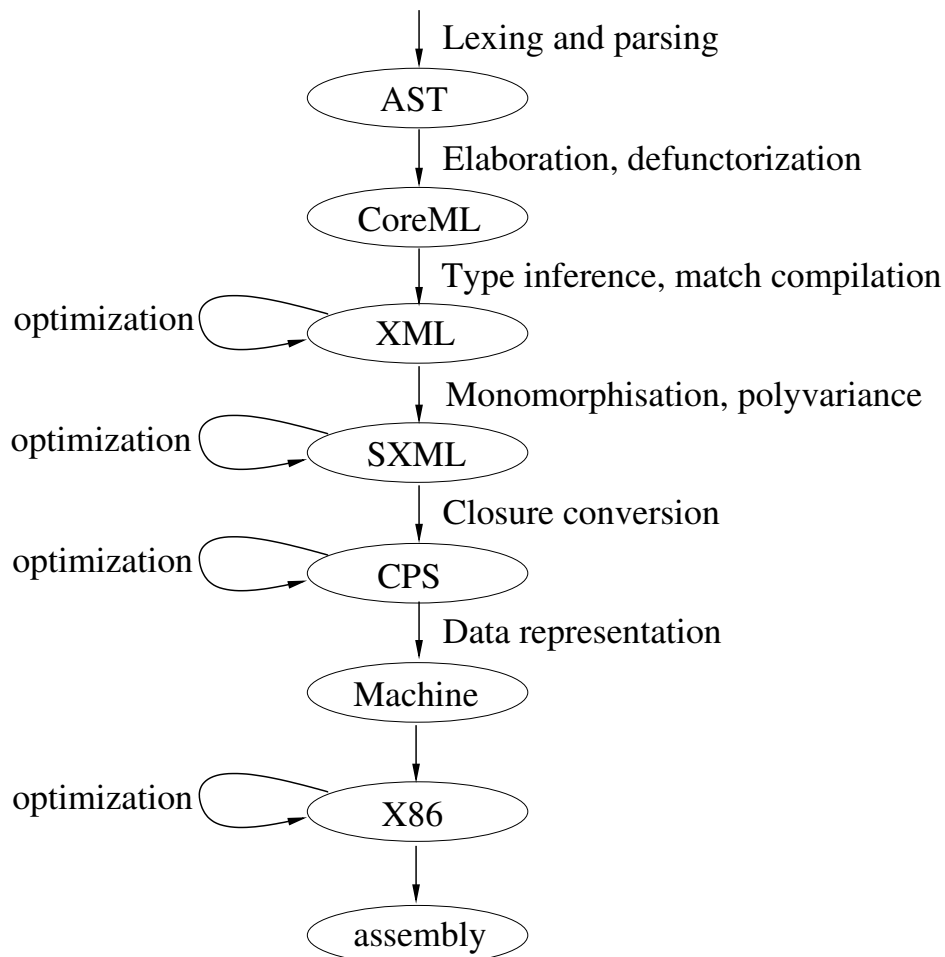


Figure 4.1: Compiler structure

4.2 Compiler Overview

Figure 4.1 shows the overall structure of the compiler. Intermediate languages (ILs) are shown in ovals. The names of compiler passes adorn arrows between ILs. In this section I give a brief description of each pass and a pointer to a later section that covers the pass in detail. Each IL also has a separate section devoted to it.

The front end (Chapter ??) takes SML source code (a complete program) and performs lexing and parsing, producing an abstract syntax tree (Chapter ??). The lexer is produced by `ml-lex[?]` and the parser is produced by `ml-yacc[?]`. The specifications for the lexer and parser were originally taken from SML/NJ109.32. The lexer is unchanged. I have substantially modified the actions in the grammar to produce my own version of abstract syntax trees (similar to, but different from SML/NJ).

Defunctionization (Chapter ??), translates abstract syntax trees to a small implicitly typed core language, called Core ML (Chapter ??). Its primary task is to eliminate all uses of the module

system (signatures, structures, functors). It does this by applying all functors and flattening all structures, moving declarations to the top level. This phase also performs precedence parsing of infix expressions and patterns (the code to do this was taken from SML/NJ). Finally, it does some amount of "macro expansion", so that the core language is smaller.

Type inference (Chapter ??) translates implicitly typed Core ML to an explicitly typed core language, XML (Chapter ??), with explicit type abstraction and application. XML is based on the language "Core-XML" described in [?]. Type inference consists of two passes. The first pass determines the binding sites of type variables that are not explicitly bound (section 4.6 of the Definition). The second pass is a pretty standard unification based Hindley-Milner type inference[?]. The type inference pass also performs overloading resolution and resolution of flexible record patterns. This pass also performs match compilation, by which I mean the translation of case statements with nested patterns to (nested) case statements with flat patterns.

Monomorphisation (Chapter ??) translates XML to its simply-typed subset, called SXML (Chapter ??), by duplicating all polymorphic functions and datatypes for each type at which they are instantiated. Monomorphisation is only possible because SML has "let-style" polymorphism, in which all uses of a polymorphic value are syntactically apparent (after functors are eliminated).

Chapter 5

Notes

This chapter contains random notes (usually old emails) on various subtle issues.

5.1 IntInf and Flattener

From: "Stephen T. Weeks" <sweeks@intertrust.com>
Date: Tue, 27 Jun 2000 18:52:19 -0700 (PDT)
To: MLton@research.nj.nec.com
Subject: safe for space ... and IntInf

Your mail also came at a fortunate time, as I was trying to track down a seg fault I was getting in the smith-normal-form regression test. For stress testing, I turned off all the cps simplify passes (except for poly equal) and ran the regressions. smith-normal-form failed with a seg fault when compiled normally, and failed with an assertion failure in `IntInf_do_neg` when compiled -g. The assertion failure was right at the beginning, checking that the frontier is in the expected place.

```
assert(frontier == (pointer)&bp->limbs[bp->card - 1]);
```

I'd been tracking this bug for a couple hours when I received your mail about the flattener. Do you see the connection? :-) As a reminder, here is the code for `bigNegate`

```
fun bigNegate (arg: bigInt): bigInt =  
  if Prim.isSmall arg  
  then let val argw = Prim.toWord arg  
       in if argw = badw  
          then negBad  
          else Prim.fromWord (Word.- (0w2, argw))  
       end  
  else Prim.~ (arg, allocate (1 + bigSize arg))
```

The problem is, when the flattener is turned off, there is an allocation in between the call to `allocate` and the `Prim.~` call. The argument tuple allocation screws everything up. So, we are

relying on the flattener for correctness of the IntInf implementation. Any ideas on how to improve the implementation to remove this reliance, or at least put an assert somewhere to avoid falling prey to this bug again?

Chapter 6

Todo

native backend vs x86 backend To unpackage debian, do
`dpkg -x ../mlton_20010806-2_i386.deb .`

Bibliography